


COMS W4701: Artificial Intelligence

Lecture 12: Active Reinforcement Learning



Instructor: Tony Dear

*Lecture materials derived from UC Berkeley's AI course at ai.berkeley.edu

Announcements

- Check HW2 grades and solutions
- Regrade requests by next Wednesday
- Regular recitations tomorrow (topic: RL)
- Lecture next Tuesday: Generalizations of RL, topics review for midterm
- Midterm review session: Example problems for midterm
- Next Wednesday 10/16, 5:30pm in CSB 451

Today

- Passive TD learning
- Multi-armed bandits
- Exploration vs exploitation
- Q-learning

Passive Reinforcement Learning

- Don't know underlying model
- Given a fixed policy, find corresponding values
- Learning is based on a set of **trials** and **observations**
- **Model-based**: Estimate and update MDP model (adaptive dynamic programming)
- **Model-free**: Do away with MDP model entirely
- Direct estimation: Accumulate estimates of expected utilities per trial
- **Temporal-difference learning**: Utility estimates persist and update each time a transition is observed based on samples

Temporal-Difference Learning


- Direct utility estimation ignores underlying model and MDP relationship
- ADP respects the MDP, but learning and updating model can get clunky
- **TD learning:** Treat each *transition* (s, a, s') as a *sample* for $V^\pi(s)$
- Keep track of V^π so far and update $V^\pi(s)$ with each transition

$$sample = R(s, \pi(s), s') + \gamma V^\pi(s')$$

α : learning rate

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$$

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$$



Example: TD Learning

States

	A	
B	C	D
	E	

Assume: $\gamma = 1$, $\alpha = 1/2$

Observed Transitions

B, east, C, -2

	0	
0	0	8
	0	

C, east, D, -2

	0	
-1	0	8
	0	

	0	
-1	3	8
	0	

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

$$V^\pi(B) \leftarrow 0.5V^\pi(B) + 0.5(-2 + 1V^\pi(C)) = 0.5(0) + 0.5(-2 + 1(0)) = -1$$



$$V^\pi(C) \leftarrow 0.5V^\pi(C) + 0.5(-2 + 1V^\pi(D)) = 0.5(0) + 0.5(-2 + 1(8)) = 3$$

Model-Based vs TD Learning

- Both preserve the underlying MDP model


$$V_{k+1}^{\pi}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$

- ADP averages over **all** possible outcomes from every state, weighted by transition probabilities
- Same as policy evaluation
- TD only averages for each observed state
- Converges more slowly and unpredictably, but simpler than ADP



Multi-Armed Bandits

- Problem: Don't know environment or optimal policy
- Still want to maximize rewards
- Tradeoff between **exploration** and **exploitation**
 - Gather more information or maximize best rewards so far?
 - How to determine when model is good enough?
- Applications: Resource allocation for maximizing productivity, clinical trials to explore different treatments, financial portfolio design



Active Reinforcement Learning

- Bandits are a special case of **active** reinforcement learning
- Unknown transitions, rewards, policy, values
- Goal: Still want to maximize expected utility
- Suppose we have value estimates from ADP or TD learning
- Then we can extract the best policy for these values

$$\hat{\pi}(s) = \operatorname{argmax}_a \sum_{s'} \hat{T}(s, a, s') [\hat{R}(s, a, s') + \gamma \hat{V}(s')]$$

- Is $\hat{\pi}$ the best policy overall? Probably not since we're relying on estimates!
- Need to **explore**; take non-optimal actions and update the model

→	→	→	+1
↓		↑	-1
→	→	↑	↓

Exploration Functions

- Simplest exploration scheme: **ϵ -greedy**
- Follow policy but perform random action with probability ϵ
- **Exploration function**: Prioritize less visited states
- Replace value estimates \hat{V} with $f(\hat{V}, N(s, a))$
- $N(s, a)$ counts number of times that (s, a) has been visited



- Suppose we want to visit all states at least N_e times
- One option: *Replace* \hat{V} with optimistic estimate R^+

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

- Alternatively: *Inflate* \hat{V} with a bonus that decreases over time $f(u, n) = u + N_e/n$


Q-Value Iteration

- We're doing an awful lot of work to **extract a new policy** every time we update our model (ADP) or value estimates (TD learning)
- How do we extract a policy? Argmax over Q-values!
- Why not just compute and keep Q-values around?

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$



$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_k(s', a')]$$



Q-Learning

- We can use samples to simulate Bellman update for Q-values instead of state values

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$

- If we take TD approach, estimate Q using running average

$$sample = r + \gamma \max_{a'} Q(s', a')$$


α : learning rate

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(sample)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha(sample - Q(s, a))$$

- Action selection if using exploration function:

$$a = \operatorname{argmax}_{a'} f(Q(s', a'), N(s', a'))$$



Example: Q-Learning

- Observed transitions: B, east, C, -2 C, north, A, -2
- Which one is exploration and which one is exploitation?
- What updates occur for Q-learning?

$$sample = r + \gamma \max_{a'} Q(s', a')$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha (sample - Q(s, a))$$

$$sample_1 = -2 + 0.8 \max(-12, -4, 4, 6) = 2.8$$

$$sample_2 = -2 + 0.8 \max(-10) = -10$$

$$Q(B, east) = 4 + 0.5(2.8 - 4) = 3.4$$

$$Q(C, north) = -12 + 0.5(-10 - (-12)) = -11$$

	<div>A</div> <div>-10</div>	
<div>0</div> <div>-2</div> <div>B</div> <div>0</div> <div>+4</div>	<div>-12</div> <div>-4</div> <div>C</div> <div>+4</div> <div>+6</div>	<div>+10</div> <div>D</div>
	<div>+4</div> <div>0</div> <div>E</div> <div>-2</div> <div>0</div>	

$$\gamma = 0.8 \quad \alpha = 0.5$$

Parameters for Q-Learning

- Q-learning updates are **off-policy**: may not reflect actions taken if exploring

$$sample = r + \gamma \max_{a'} Q(s', a')$$

- Can still converge to optimal policy!
- **Learning rate α** : Convergence guaranteed if α decreases to 0 over time
 - In practice, a constant rate, e.g. $\alpha = 0.1$, is sufficient
- **Exploration rate ε** : Can be constant, can decrease over time depending on context
- **Discount factor γ** : Determines significance of future rewards to agent
 - In practice, learning is faster if starting with lower γ and then increasing over time
- **Initial conditions Q_0** : Inflated initial values can encourage exploration

Summary: MDPs and RL

Known MDP: Offline Dynamic Programming

Goal

Evaluate fixed policy π

Find optimal $\pi^*, V^*, (Q^*)$

Technique

Policy evaluation

Value / policy iteration

Unknown MDP: Model-Based

Goal

Evaluate fixed policy π

Find optimal $\pi^*, V^*, (Q^*)$

Technique

ADP / policy evaluation

ADP / policy exploration

Unknown MDP: Model-Free

Goal

Evaluate fixed policy π

Find optimal $\pi^*, V^*, (Q^*)$

Technique

TD Value Learning

Q-learning