


# COMS W4701: Artificial Intelligence

## Lecture 17: Inference in Hidden Markov Models




Instructor: Tony Dear

\*Lecture materials derived from UC Berkeley's AI course at [ai.berkeley.edu](https://ai.berkeley.edu)

# HMM Conditional Independences

- Markov chain independences:

$$X_t \perp\!\!\!\perp X_1, \dots, X_{t-2} \mid X_{t-1} \quad \begin{matrix} P(X_t \mid X_{t-1}) \\ P(E_t \mid X_t) \end{matrix}$$



- A state is conditionally independent of past states and evidence given preceding state:


$$X_t \perp\!\!\!\perp X_1, E_1, \dots, X_{t-2}, E_{t-2}, E_{t-1} \mid X_{t-1}$$

- An emission is conditionally independence of past states and evidence given current state:

$$E_t \perp\!\!\!\perp X_1, E_1, \dots, X_{t-2}, E_{t-2}, X_{t-1}, E_{t-1} \mid X_t$$



# HMM Joint Distribution



- General joint distribution:

$$P(X_1, E_1, \dots, X_T, E_T) = P(X_1)P(E_1|X_1) \prod_{t=2}^T P(X_t|X_{t-1})P(E_t|X_t)$$

- Marginal distributions can be found by summing out RVs
- For certain computations we don't even need the entire joint distribution!

# Applications of HMMs

---

- **Speech recognition**
  - Observations: Acoustic signals / waveforms
  - States: Positions in words
- **Machine translation**
  - Observations: Words to be translated
  - States: Translation options
- **Robot tracking**
  - Observations: Range readings
  - States: Positions on a map

# Today

---

- Inference tasks in hidden Markov models
- State estimation (filtering): Forward algorithm
- Most likely explanation: Viterbi algorithm


# HMMs and Inference

---

- We are generally interested in hidden states  $X$  given *observed* evidence  $e$
- **Filtering** (state estimation): Find  $P(X_t \mid e_{1:t})$ 
  - What is the hidden state, given *all evidence to date*?
- **Most likely explanation**: Find  $\operatorname{argmax}_{x_{1:t}} P(X_{1:t} \mid e_{1:t})$ 
  - What is the *sequence* of hidden states that best explains the observed evidence?
- **Smoothing**: Find  $P(X_k \mid e_{1:t})$ , for  $1 \leq k < t$ 
  - Use both past and future evidence to *smooth* prediction of a state

# Inference: State Estimation


- Idea: Track a *belief state* over time:  $P(X_t | e_{1:t})$
- We want to compute this recursively (constant time)



- For each timestep, we update our belief as follows:
- *Elapse* time: Follow the state transition model (same as Markov chains)
- *Observe* evidence: Follow the emissions model to update belief

# Example: Robot Localization

Example from  
Michael Pfeiffer




Prob

0

1


$t=0$




- Hidden state: Robot's true location
- **Motion** (transition) model
  - Move in *intended* direction with larger probability)
- **Sensor** (emissions) model
  - Wall or no wall in each cardinal direction, noisy readings



# Example: Robot Localization




$t=1$




- One observation, before moving:
- Can narrow down to non-corner states at top and bottom
- Note the less likely (but prob non-zero) states in the middle!

# Example: Robot Localization




t=2




- Move and sense two walls again
- More likely for robot to be in the middle darker locations
- Light grey on left: More likely that we moved than not

# Example: Robot Localization




t=3




- Robot continues updating its belief state of where it is...

# Example: Robot Localization




$t=4$




- Robot continues updating its belief state of where it is...

# Example: Robot Localization



t=5



- Robot is now very confident in its belief about its current location!

# Normalization

- We want to find  $P(X_t \mid e_{1:t})$ —use def of conditional probability:

$$P(X_t \mid e_{1:t}) = \frac{P(X_t, e_{1:t})}{P(e_{1:t})}$$

- Denominator corresponds to *observed* random variables
- We can compute this, but this is also just a constant (why?)
- Instead of trying to compute  $P(e_{1:t})$ , we can simply *normalize*  $P(X_t, e_{1:t})$

$$P(X_t \mid e_{1:t}) = \alpha P(X_t, e_{1:t}) \propto_{X_t} P(X_t, e_{1:t})$$

# Forward Algorithm

- Let's suppose we have  $f_t = P(X_t | e_{1:t})$

- Elapse time:**

$$\sum_{x_t} P(x_t | e_{1:t}) P(X_{t+1} | x_t, e_{1:t}) = \sum_{x_t} P(x_t, X_{t+1} | e_{1:t})$$

Transition

HMM says that state is independent of all prior stuff in the presence of an immediate state

Conditional independence

Joint probability for  $x_t, x_{t+1}$

$$= P(X_{t+1} | e_{1:t})$$

→  $f_t \cdot P(X_{t+1} | X_t) = f'_{t+1}$

- Observe evidence:**

$$P(x_{t+1} | e_{1:t}) P(e_{t+1} | x_{t+1}, e_{1:t}) = P(x_{t+1}, e_{t+1} | e_{1:t})$$

Emission

Conditional independence

→  $f'_{t+1} * P(e_{t+1} | X_{t+1}) \propto_{X_{t+1}} f_{t+1}$

Pointwise multiply      Normalize

$\propto_{X_{t+1}} P(x_{t+1} | e_{1:t+1})$   
Normalize


# Forward Algorithm

- Updates with *constant* time and space complexity despite unbounded sequence of observations
- Base case: Observe evidence for initial distribution:  $f_1 \propto_{X_1} f_0 * P(e_1 | X_1)$

- Elapse time *increases uncertainty*

$$f'_{t+1} = f_t \cdot P(X_{t+1} | X_t)$$


$$(p_1 \quad \cdots \quad p_n) \begin{pmatrix} p_{1|1} & \cdots & p_{n|1} \\ \vdots & \ddots & \vdots \\ p_{1|n} & \cdots & p_{n|n} \end{pmatrix}$$



- Observation reweights beliefs, *decreases uncertainty*

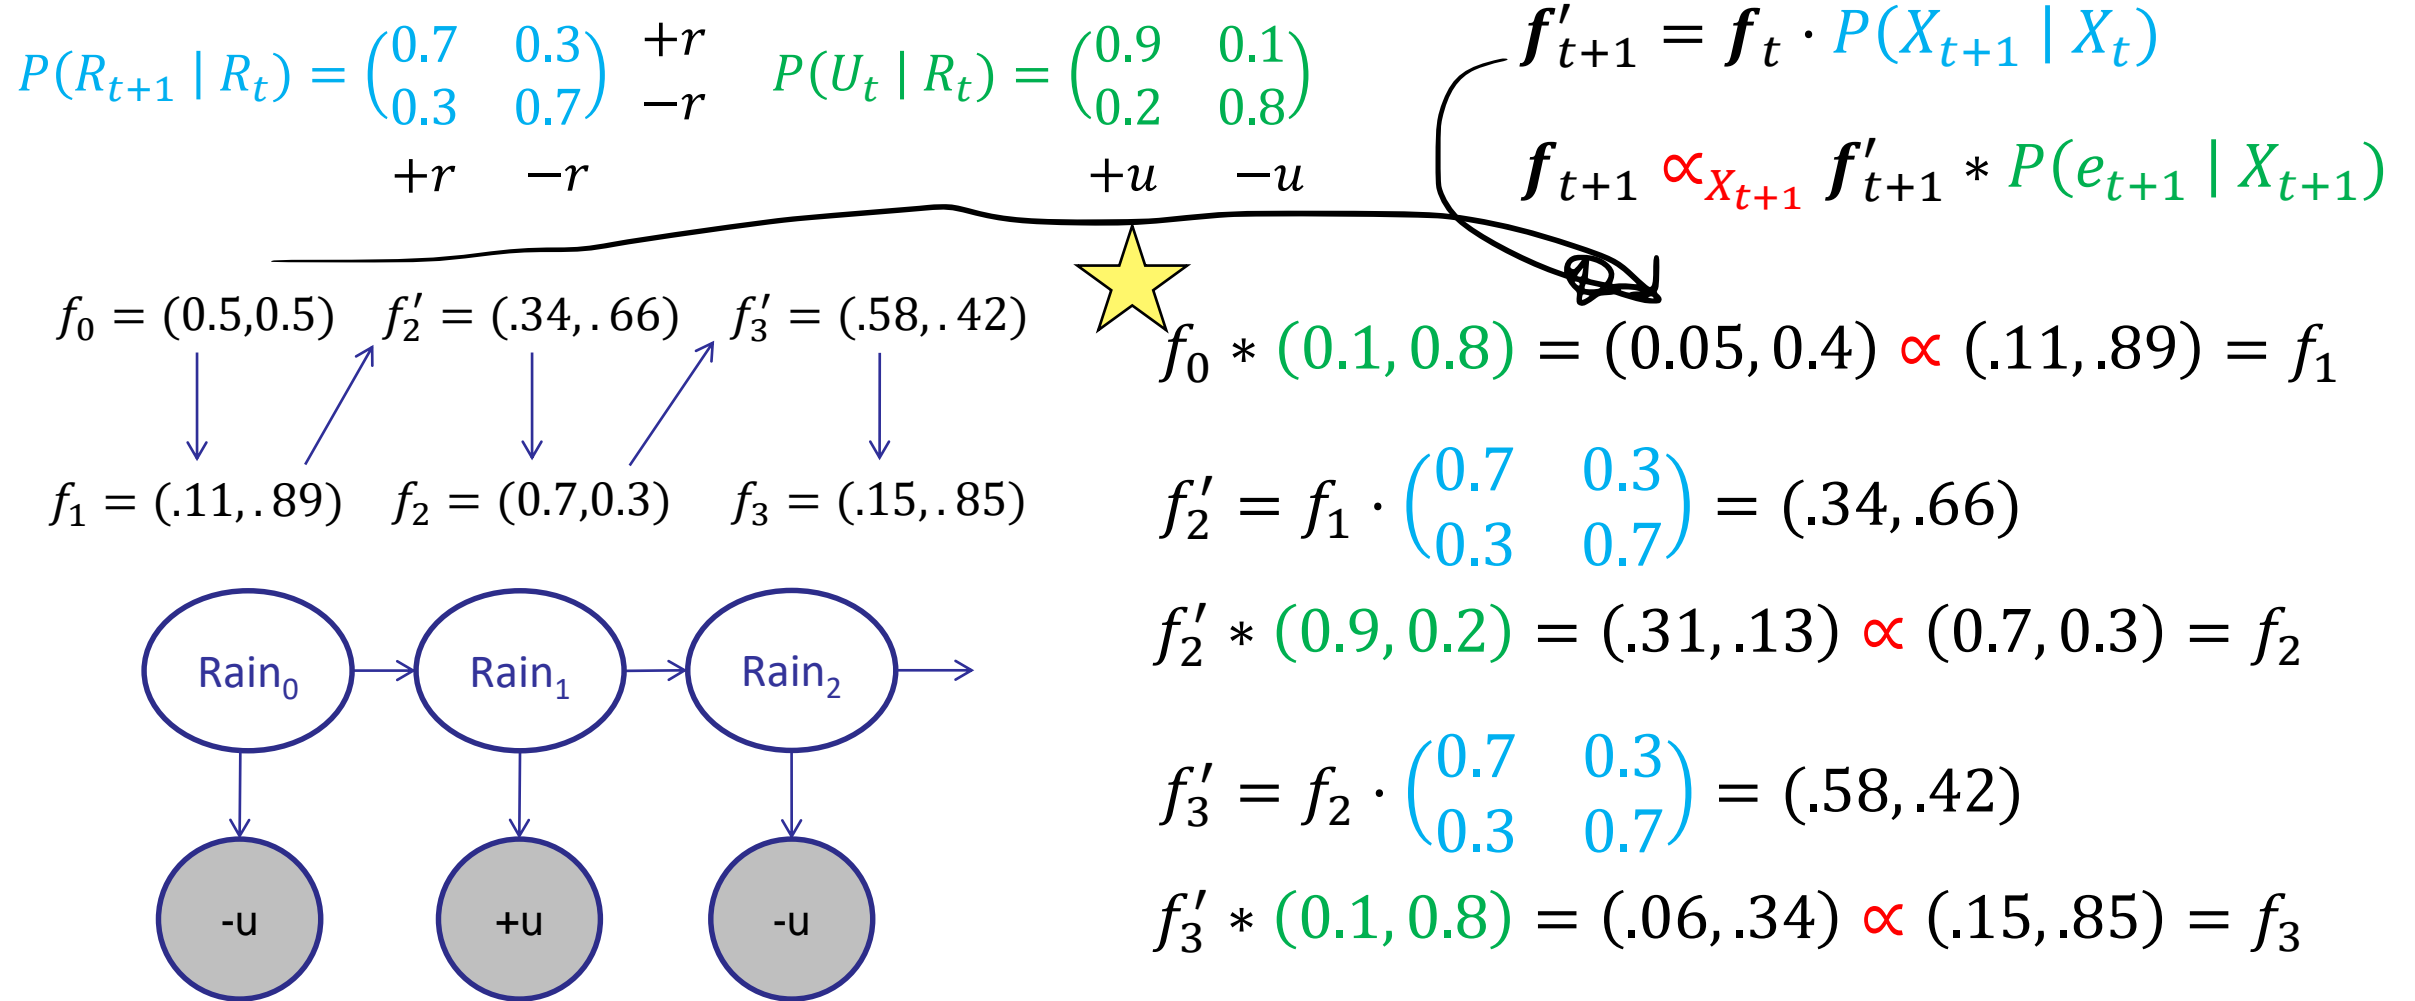
$$f_{t+1} \propto_{X_{t+1}} f'_{t+1} * P(e_{t+1} | X_{t+1})$$

$$(p'_1 \quad \cdots \quad p'_n) * \begin{pmatrix} \ddots & p_{e|1} & \ddots \\ \ddots & \vdots & \ddots \\ \ddots & p_{e|n} & \ddots \end{pmatrix}$$






# Example: Weather HMM



# Inference: Most Likely Sequence



# Most Likely Sequence

- What is the most likely *sequence* of states given a *sequence* of evidence?

$$\operatorname{argmax}_{x_{1:t}} P(X_{1:t} \mid e_{1:t})$$

- Equivalently, we can argmax the **joint probability**  $P(X_{1:t} \mid e_{1:t}) \propto P(X_{1:t}, e_{1:t})$

- We **cannot** just run forward algorithm for each state and find argmax!

- Most likely individual states may differ from that of the most likely sequence

$$\operatorname{argmax} P(X_1) = +x$$


$$\operatorname{argmax} P(X_2) = -x$$

$$\operatorname{argmax} P(X_3) = -x$$

$X_1$	$X_2$	$X_3$	$P(X_1, X_2)$
$+x$	$+x$	$+x$	0.05
$+x$	$+x$	$-x$	0.1
$+x$	$-x$	$+x$	0.3
$+x$	$-x$	$-x$	0.15
$-x$	$+x$	$+x$	0
$-x$	$+x$	$-x$	0.2
$-x$	$-x$	$+x$	0.05
$-x$	$-x$	$-x$	0.15

$$\operatorname{argmax} P(X_1, X_2, X_3) = (+x, -x, +x)$$

# State Trellis Diagram



- A state sequence is a *path* through a **state trellis diagram**
- Each arc is a transition  $x_{t-1} \rightarrow x_t$  with weight  $P(e_t | x_t)P(x_t | x_{t-1})$
- A state sequence is also a specific *event* of joint state values
- Maximizing the joint probability = maximizing the *product* of arc weights along a path
- Idea: Best path to state  $x_t$  *includes* best path to state  $x_{t-1}$ , followed by a transition
- *Recursively* compute best paths by recording max joint probabilities so far

# Viterbi Algorithm

- $\mathbf{m}_t = \max_{x_1 \dots x_{t-1}} P(x_{1:t-1}, \mathbf{X}_t, e_{1:t}) \propto \max_{x_1 \dots x_{t-1}} P(x_{1:t-1}, \mathbf{X}_t \mid e_{1:t})$  is a distribution over  $X_t$
- Each  $\mathbf{m}_t(x_t)$  is a joint probability of most likely sequence up to  $x_t$
- Then  $\mathbf{m}_{t+1}(x_{t+1})$  “concatenates”  $\mathbf{m}_t$  with a state value  $x_t$ :

$$\begin{aligned}
 \mathbf{m}_{t+1}(x_{t+1}) &= \max_{x_1 \dots x_t} P(x_{1:t}, x_{t+1}, e_{1:t+1}) \\
 &= \max_{x_1 \dots x_t} P(x_{1:t-1}, x_t, e_{1:t}) P(x_{t+1} \mid x_t, \cancel{x_{1:t-1}}, \cancel{e_{1:t}}) P(e_{t+1} \mid x_{t+1}, \cancel{x_t}, \cancel{x_{1:t-1}}, \cancel{e_{1:t}}) \\
 &= \max_{x_1 \dots x_t} P(x_{1:t-1}, x_t, e_{1:t}) P(x_{t+1} \mid x_t) P(e_{t+1} \mid x_{t+1}) \\
 &= \max_{x_t} P(e_{t+1} \mid x_{t+1}) P(x_{t+1} \mid x_t) \max_{x_1 \dots x_{t-1}} P(x_{1:t-1}, x_t, e_{1:t}) \\
 &= \underbrace{P(e_{t+1} \mid x_{t+1})}_{\text{Emission}} \max_{x_t} \underbrace{P(x_{t+1} \mid x_t)}_{\text{Transition}} \mathbf{m}_t(x_t)
 \end{aligned}$$

Conditional  
independence

Conditional  
independence

Same as forward algorithm  
but replace sum with max!

# Viterbi Algorithm

- **Elapse time:** Each value of  $\mathbf{m}'_{t+1}$  maxes over *pointwise product* between  $\mathbf{m}_t$  and corresponding *column* of transition matrix

$$\mathbf{m}'_{t+1}(x_{t+1}) = \max(\mathbf{m}_t * P(x_{t+1} | X_t))$$

$$(p_1 \cdots p_n) \begin{pmatrix} p_{1|1} & \cdots & p_{n|1} \\ \vdots & \ddots & \vdots \\ p_{1|n} & \cdots & p_{n|n} \end{pmatrix} \begin{matrix} * x_{t+1} \\ \boxed{\begin{matrix} \cdots \\ \vdots \\ \cdots \end{matrix}} \\ \text{max} \end{matrix}$$

- Since we want a *sequence of states*, we also need a pointer to best *parent* of each  $x_{t+1}$ :  $Pointer_{t+1}(x_{t+1}) = \operatorname{argmax}_{x_t}(\mathbf{m}_t * P(x_{t+1} | X_t))$
- **Observe evidence:** No need to normalize (why?)  $\mathbf{m}_{t+1} = \mathbf{m}'_{t+1} * P(e_{t+1} | X_{t+1})$
- *Backward pass:* Starting with  $Pointer_T(\max \mathbf{m}_T)$ , follow pointers backwards to  $x_1$  to extract most likely sequence of states


# Example: Weather HMM

$$P(R_{t+1} | R_t) = \begin{pmatrix} 0.6 & 0.4 \\ 0.3 & 0.7 \end{pmatrix} \begin{matrix} +r \\ -r \end{matrix}$$

$$P(U_t | R_t) = \begin{pmatrix} 0.9 & 0.1 \\ 0.2 & 0.8 \end{pmatrix} \begin{matrix} +r \\ -r \end{matrix}$$

Observed evidence:

$$\begin{matrix} e_1 = -u \\ e_2 = +u \\ e_3 = -u \end{matrix}$$



Backward pointers:

$\text{argmax}_{x_t} \mathbf{m}_{t+1}(x_{t+1})$

Diagram showing backward pointers for the most likely sequence:

- From  $m_3$  to  $m_2$ :  $[-r]$
- From  $m_2$  to  $m_1$ :  $[+r]$

$$m_0 = (0.5, 0.5)$$

$$m_1 = m_0 * (0.1, 0.8) = (0.05, 0.4)$$

$$m'_2(+r) = \max((0.05, 0.4) * (0.6, 0.3)) = .12$$

$$m'_2(-r) = \max((0.05, 0.4) * (0.4, 0.7)) = .28$$

$$m_2 = (.12, .28) * (0.9, 0.2) = (.108, .056)$$

$$m'_3(+r) = \max((.11, .06) * (0.6, 0.3)) = .065$$

$$m'_3(-r) = \max((.11, .06) * (0.4, 0.7)) = .043$$

$$m_3 = (.065, .043) * (0.1, 0.8) = (.006, .035)$$

Most likely sequence:  $(-r, +r, -r)$

# Inference Applications

---

- Forward algorithm has linear time and constant space complexity
- Viterbi algorithm has linear time *and* linear space complexity
- Both are heavily used in digital signals (cellular, satellite, LAN, etc.), speech recognition (audio to text), bioinformatics (gene decoding), finance (stock, asset trends)
- Forward algorithm can be combined with a *backward algorithm* to perform smoothing
- Smoothing can then be used to *learn* unknown HMM model parameters using the **Baum-Welch algorithm**